

ROIperf: A Framework to Rapidly Validate Workload Sampling Methodologies

Alen Sabu¹, Harish Patil², Wim Heirman², Trevor E. Carlson¹

¹National University of Singapore

²Intel Corporation

Workshop on Computer Architecture Modeling and Simulation (CAMS 2023)

October 28, 2023




NUS
National University
of Singapore



Microarchitectural simulation is slow

Microarchitectural simulation is slow




Solution: Simulate regions of interest (sample)

Simulation wall times of SPEC CPU2017 benchmarks (8 threads) using train inputs on Sniper

Sampled Simulation

ROIperf




Solution: Simulate regions of interest (sample)

Simulation wall times of SPEC CPU2017 benchmarks (8 threads) using train inputs on Sniper

Sampled Simulation

ROIperf



Simulation wall times of SPEC CPU2017 benchmarks (8 threads) using train inputs on Sniper


Application	Train	Ref
603.bwaves_s.1	33.33	1.01
603.bwaves_s.2	32.79	1.03
607.cactuBSSN_s.1	26.81	0.45
619.lbm_s.1	4.86	0.65
621.wrf_s.1	9.28	0.47
627.cam4_s.1	4.78	0.23
628.pop2_s.1	6.27	0.46
638.imagick_s.1	25.93	0.13
644.nab_s.1	21.15	0.32
644.nab_s.2	9.74	N/A
649.fotonik3d_s.1	12.93	1.55
654.roms_s.1	3.98	0.71
657.xz_s.1	21.43	0.74
657.xz_s.2	42.55	1.26

Fraction of regions to be simulated in detail for SPEC CPU2017 benchmarks using 8 threads

Program executions are structured as phases


Selection of Regions of Interest

Program executions are structured as phases

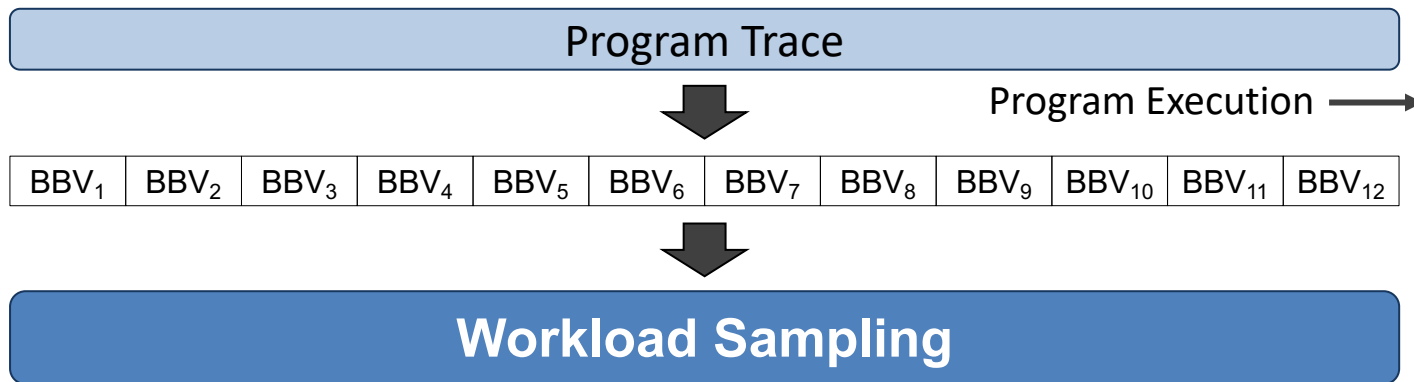


Selection of Regions of Interest

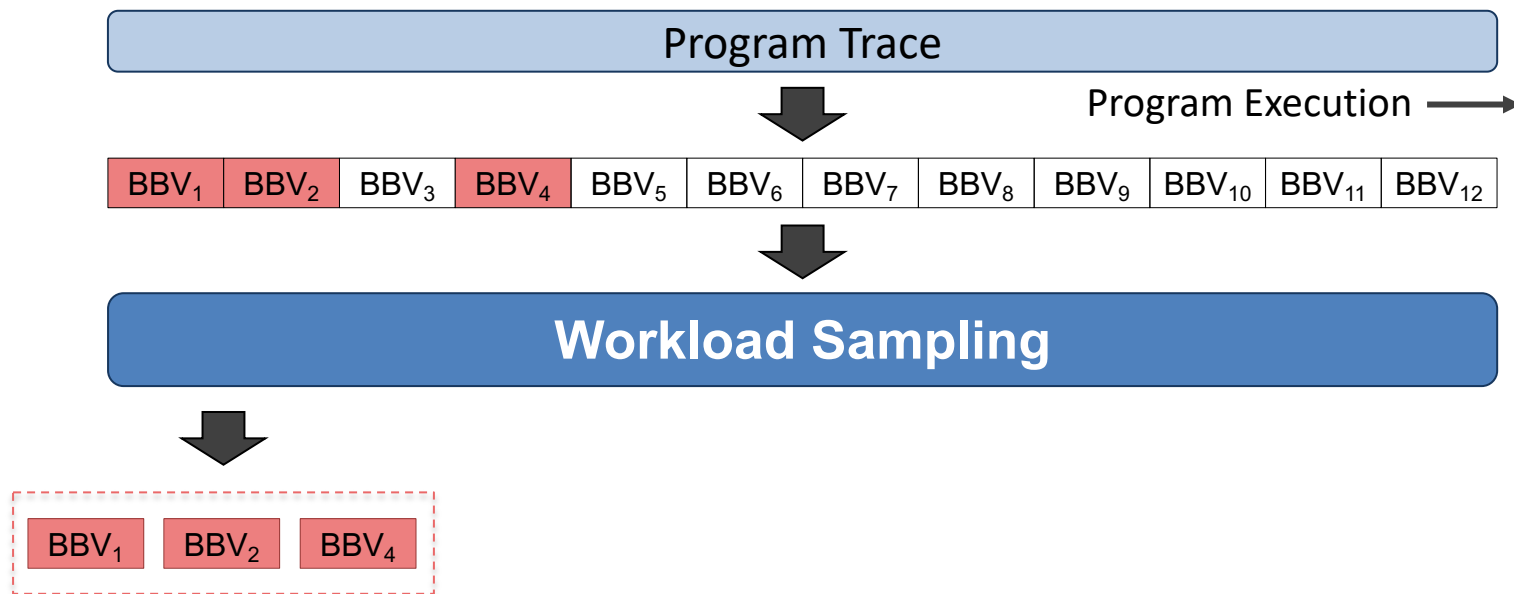
Program executions are structured as phases



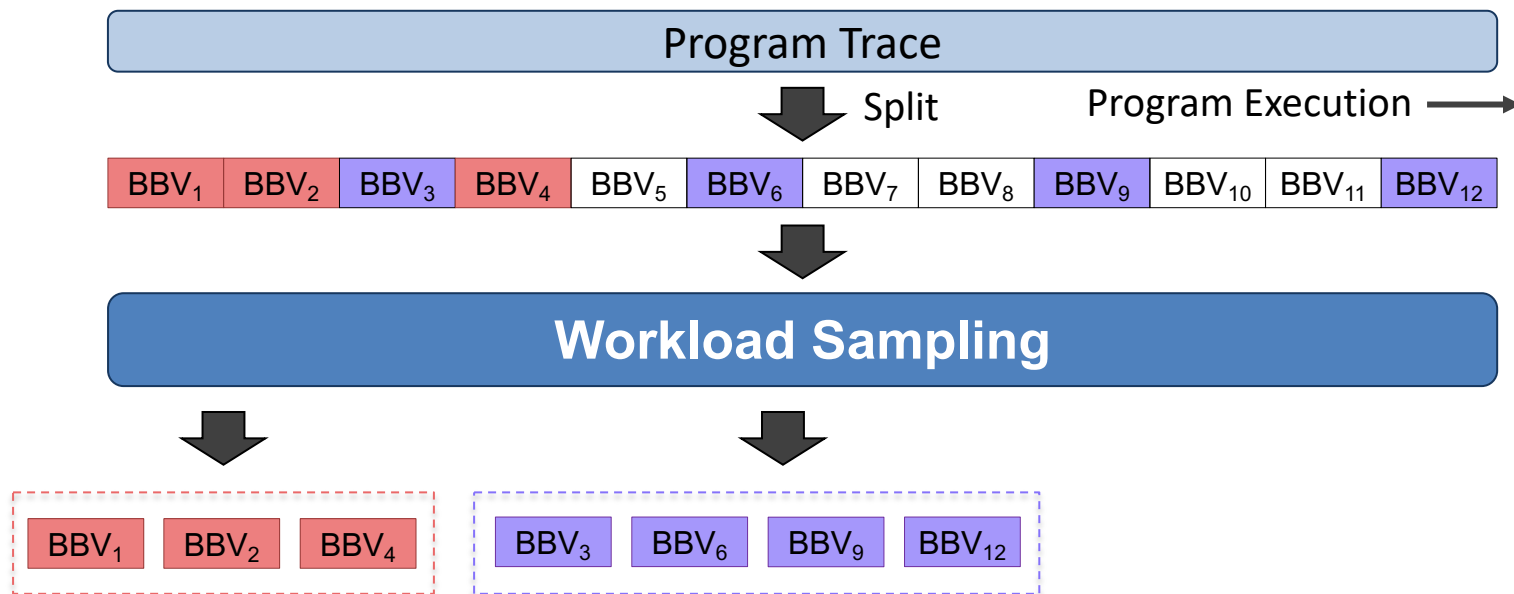
Selection of Regions of Interest



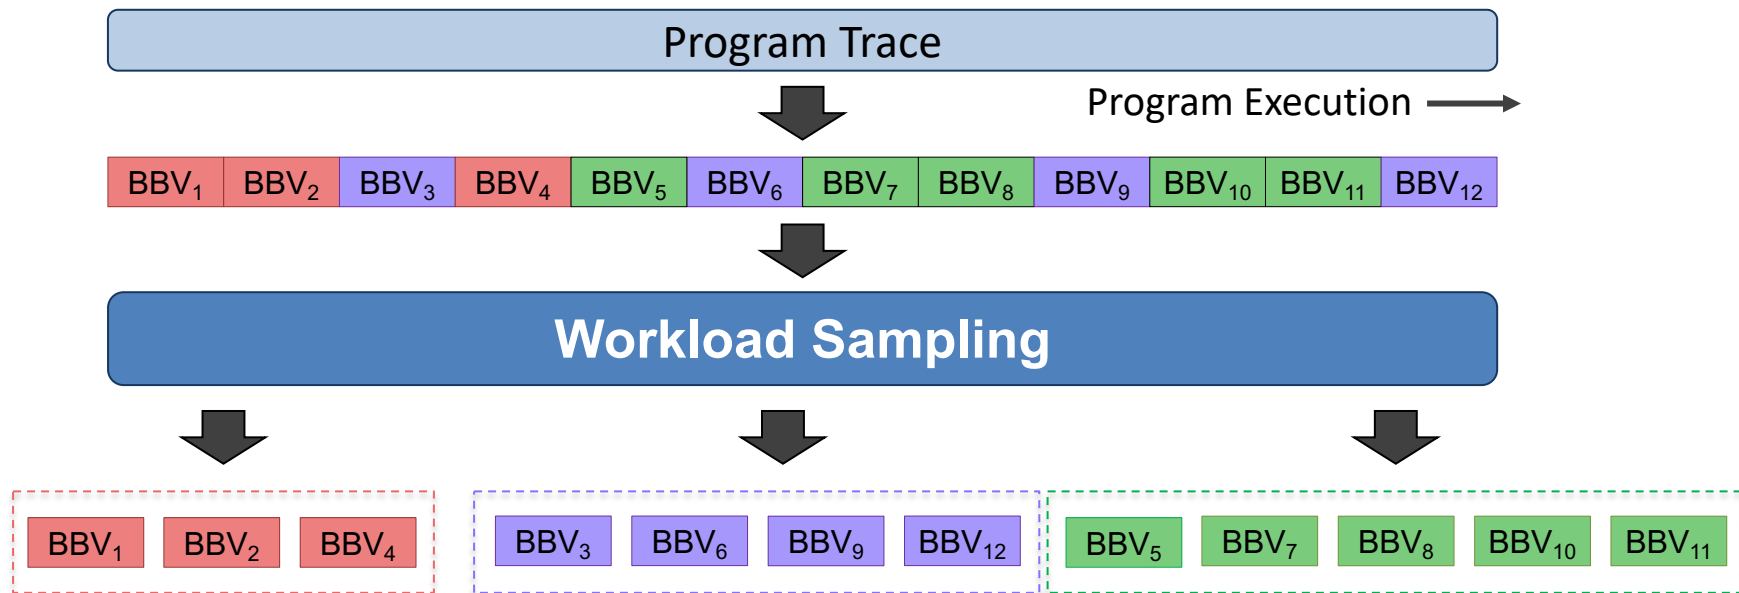
Selection of Regions of Interest



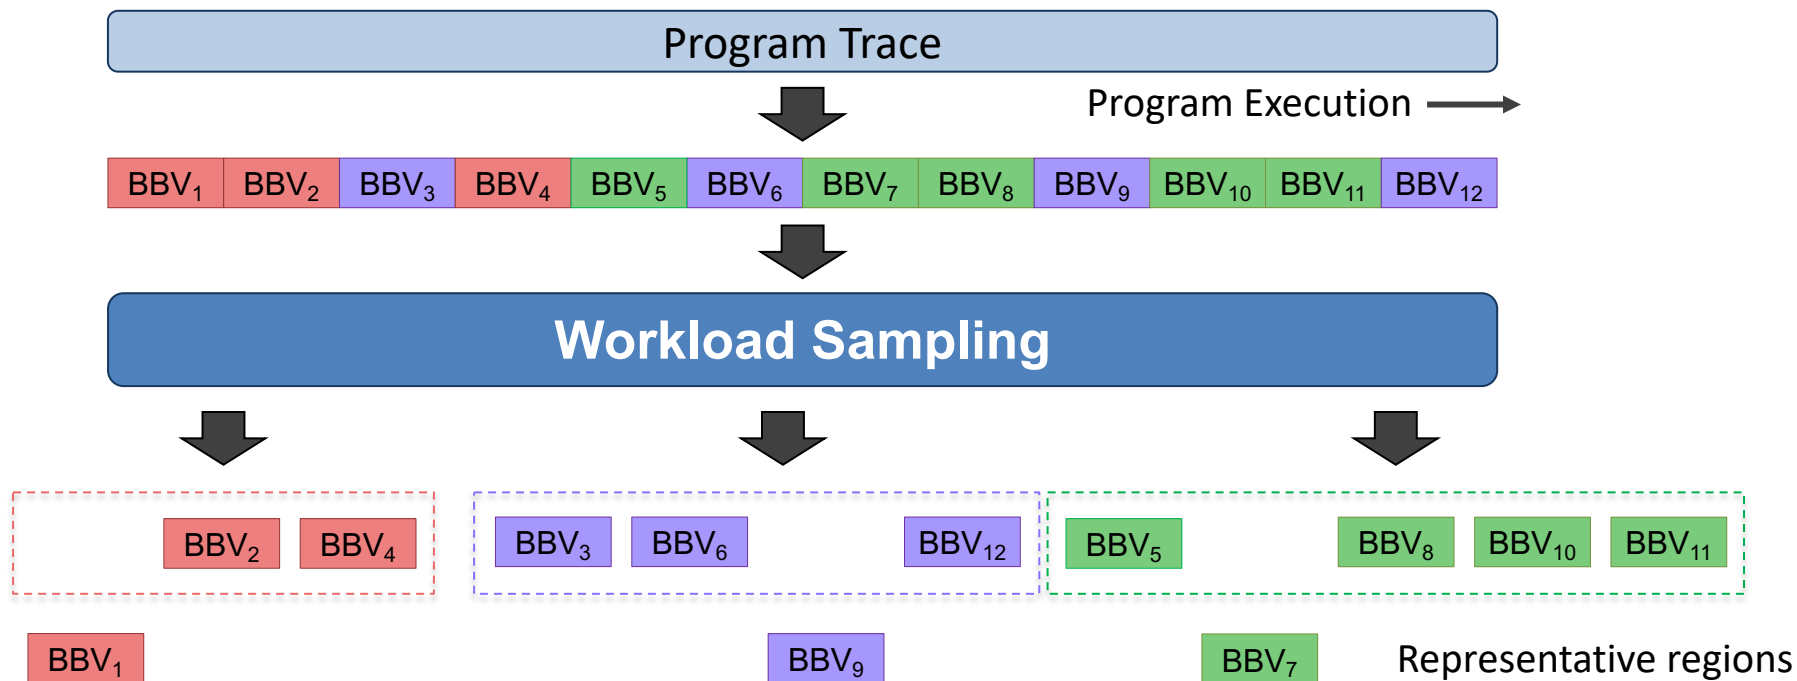
Selection of Regions of Interest



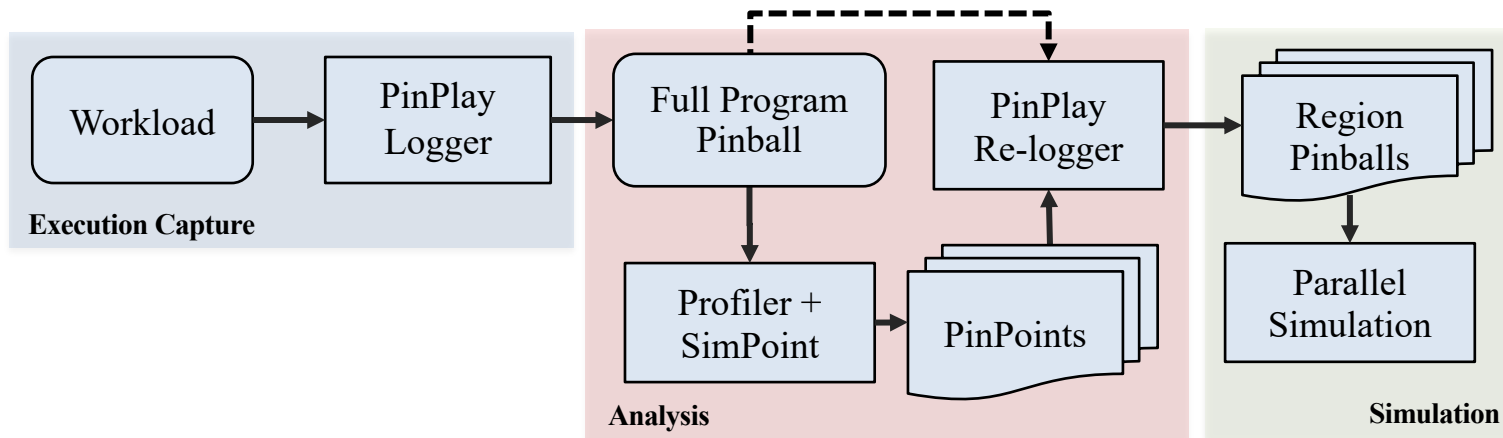
Selection of Regions of Interest



Selection of Regions of Interest

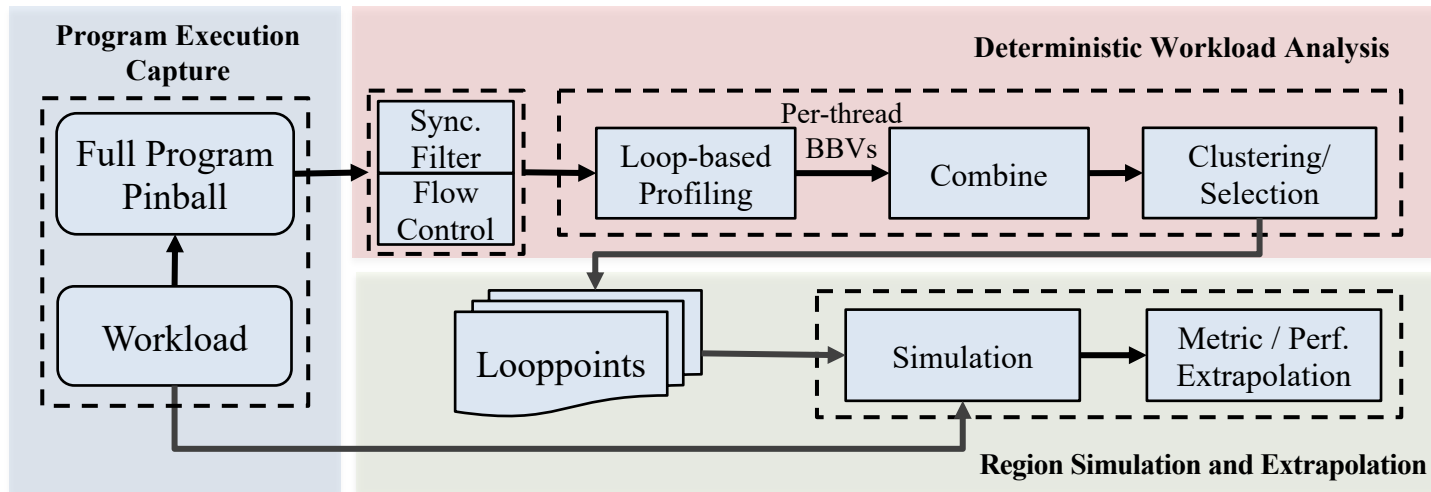


Region Selection: SimPoint



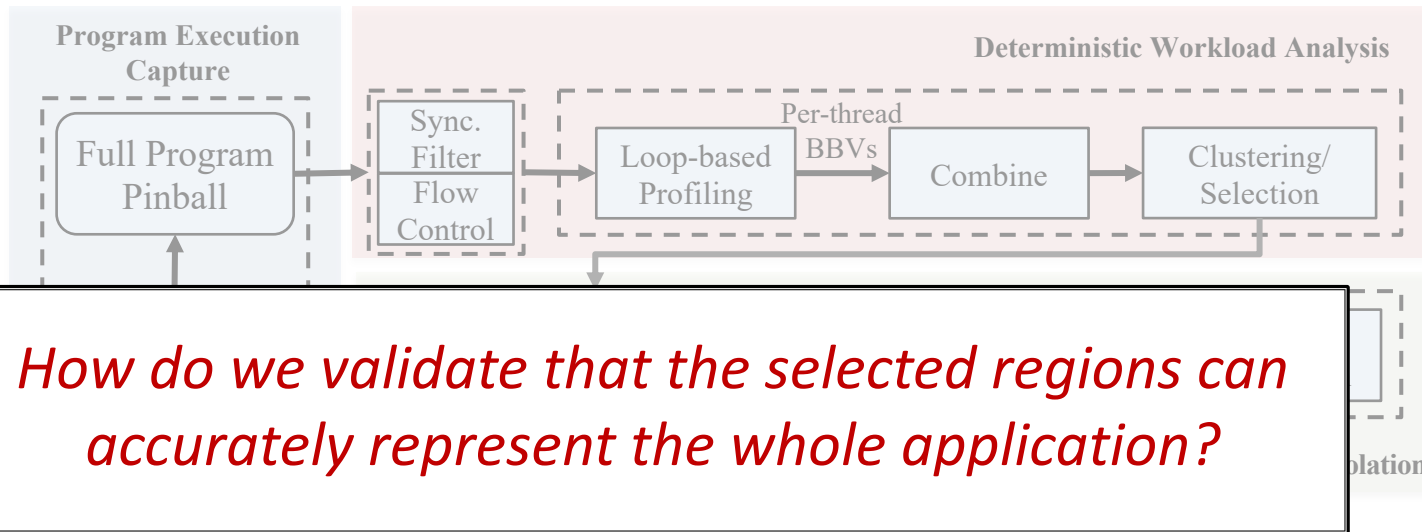
Applicable to Single-threaded Workloads

Region Selection: LoopPoint



Applicable to Multi-threaded Workloads

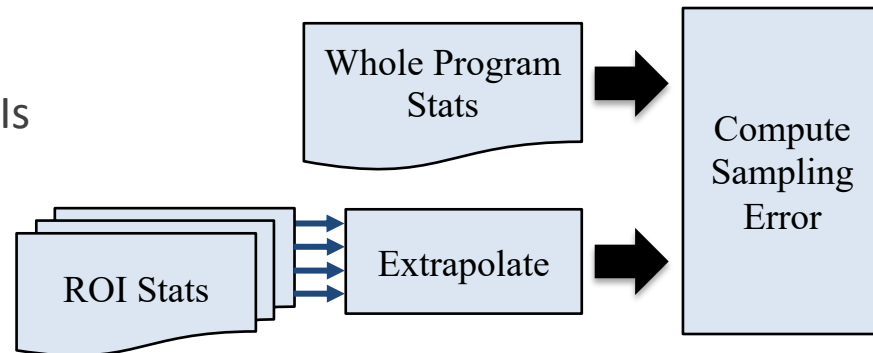
Region Selection: LoopPoint



How do we validate that the selected regions can accurately represent the whole application?

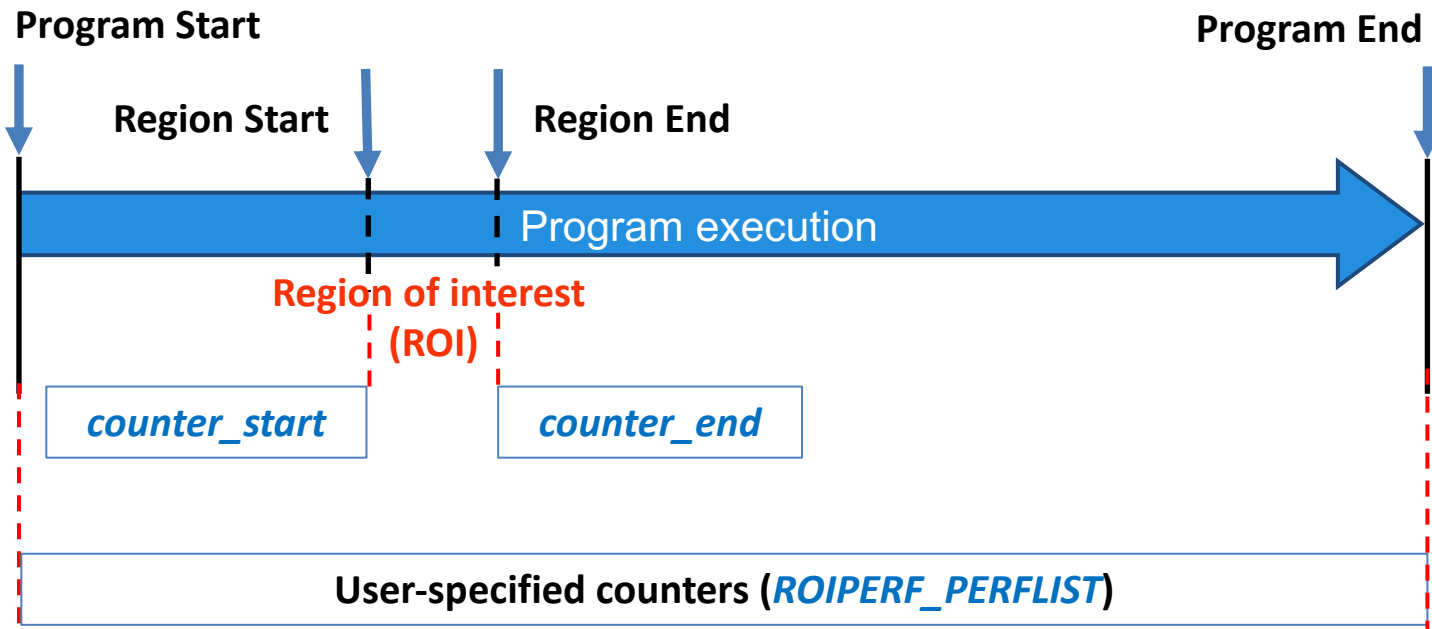
Applicable to Multi-threaded Workloads

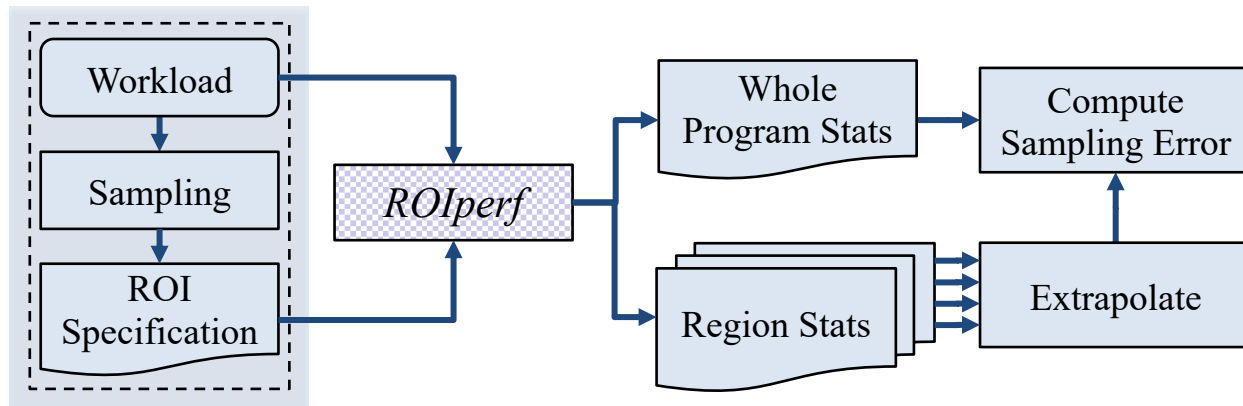
- Estimate the *representativeness* of ROI selection
 - Sampling Error
 - Performance projection using ROIs
- Performance measurement
 - Simulation
 - H/w perf counters using *ROIperf*
- Iterative fine tuning of sampling methodologies



ROIperf Workflow

ROIperf

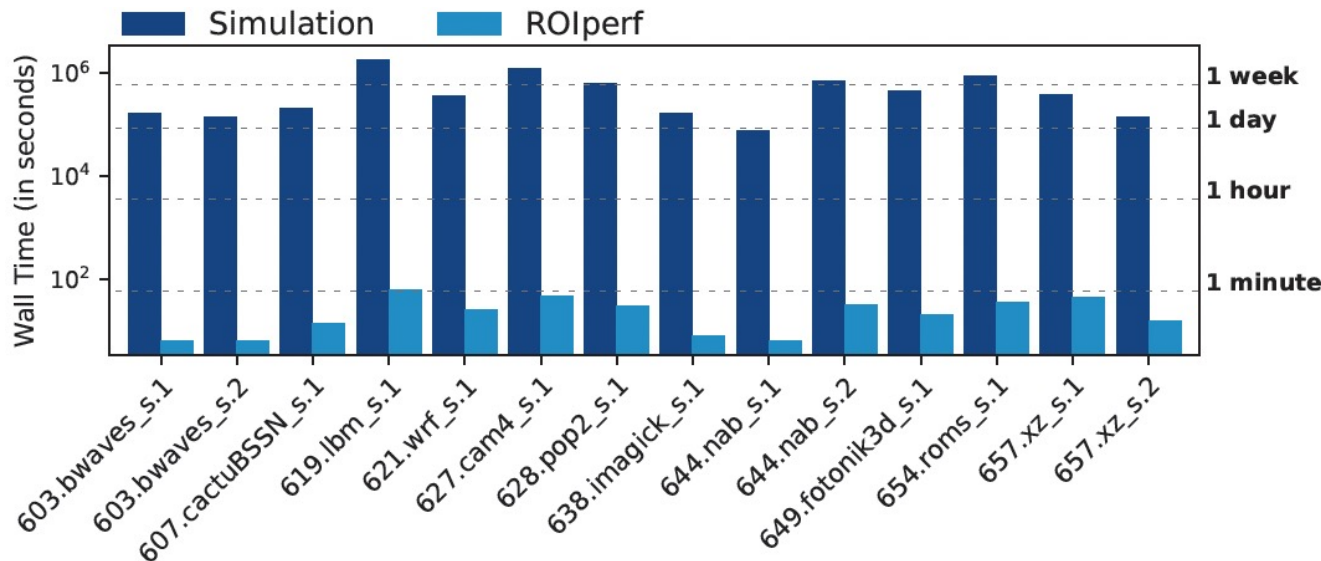




- Pin probe tool – minimize perturbation
- Outputs hardware perf counter values at start and end of ROI (Ex. RDTSC)
- Ref: `/usr/include/linux/perf_event.h`, `perf_event_open()` syscall

Wall Time Comparison for Validation

ROIperf




Simulation and ROIperf wall times of SPEC CPU2017 benchmarks (8 threads) using train inputs on Sniper and Skylake machine, respectively.


$$\text{prediction error}\% = \left(1 - \frac{\text{extrapolated metric}}{\text{full run metric}}\right) \times 100$$

$$\text{prediction error}\% = \left(1 - \frac{\text{extrapolated metric}}{\text{full run metric}}\right) \times 100$$

Single-threaded¹ (SimPoint)




Multi-threaded² (LoopPoint)



SPEC CPU2017 benchmarks with training inputs

Missing values caused by reproducibility issues!



SPEC CPU2017 benchmarks with reference inputs

- ROIperf provides a framework to **rapidly validate ROIs** selected by sampling methodologies on real hardware
- Using ROIperf, **tuning a sampling methodology** for a workload can be quickly performed
- We show the efficacy of ***SimPoint*** for single-threaded and ***LoopPoint*** for multithreaded workloads, respectively.

ROIperf: A Framework to Rapidly Validate Workload Sampling Methodologies

Alen Sabu¹, Harish Patil², Wim Heirman², Trevor E. Carlson¹

¹National University of Singapore

²Intel Corporation

Workshop on Computer Architecture Modeling and Simulation (CAMS 2023)

October 28, 2023



NUS
National University
of Singapore

